A natural encoding of genetic variation in a Burrows-Wheeler Transform to enable mapping and genome inference

Sorina Maciuca, Carlos delOjo Elias, Gil McVean, and Zamin Iqbal

Wellcome Trust Centre for Human Genetics,
University of Oxford, Roosevelt Drive, Oxford OX3 7BN, UK
{sorina.maciuca,gil.mcvean,zamin.iqbal}@well.ox.ac.uk
carlos.delojoelias@ndm.ox.ac.uk

Abstract. We show how positional markers can be used to encode genetic variation within a Burrows Wheeler Transform (BWT), and use this to construct a generalisation of the traditional "reference genome", incorporating known variation within a species. Our goal is to support the inference of the closest mosaic of previously known sequences to the genome(s) under analysis.

Our scheme results in an increased alphabet size, and by using a wavelet tree encoding of the BWT we reduce the performance impact on rank operations. We give a specialised form of the backward search that allows variation-aware exact matching. We implement this, and demonstrate the cost of constructing an index of the whole human genome with 8 million genetic variants is 25GB of RAM. We also show that inferring a closer reference can close large kilobase-scale coverage gaps in *P. falciparum*.

Keywords: pan-genome, Burrows-Wheeler Transform, FM index, genome

1 Introduction

Next-generation genome sequencing involves breaking DNA into fragments, identifying substrings (called "reads"), assembling them and then inferring properties of the genome. Recently, it has become possible to study within-species genetic variation on a large scale [6, 7], where the dominant approach is to match substrings to the canonical "reference genome" which is constructed from an arbitrary individual. This problem ("mapping") has been heavily studied (see [5]) and the Burrows Wheeler Transform (BWT) [2] underlies the two dominant mappers [3, 4]. Mapping reads to a reference genome is a very effective way of detecting genetic variation caused by single character changes (SNPs - single nucleotide polymorphisms). However, this method becomes less effective the further the genome differs from the reference. Sequences that diverge from the reference genome are difficult to map and tend to produce mapping failure. This is an important problem to address since, in many organisms, biologically relevant genomic regions are highly diverse.

For a given species, our goal is to build a compact representation of the genomes of N individuals that supports and facilitates the following inference: we take as input sequence data from a new sample, an estimate of how many genomes the sample contains and their relative proportions - e.g. a normal human sample would contain 2 genomes in a 1:1 ratio, a bacterial isolate would contain 1 genome and a malaria sample might contain 3 genomes in the ratio 10:3:1. We would then infer the sequence of the underlying genomes. We call such a data structure a Population Reference Genome (PRG). In this paper we describe a method for encoding genetic variation designed to enable this approach.

Genomes evolve mainly via two processes - mutation (changing a single character or, less frequently, inserting or deleting a few) and recombination (either two chromosomes exchange a chunk of DNA, or one chromosome copies a chunk from another). Thus once we have seen many genomes of a given species, a new genome is likely to look like a mosaic of genomes we have seen before. If we can infer a close mosaic, we have found a "personalised reference genome", and reads are more likely to match exactly. This approach was first described in [8], applied to the human MHC region. However their implementation was quite specific to the region and would not scale to the whole genome. Valenzuela *et al.* [1] have also espoused a find-the-closest-reference approach.

Other "reference graph" methods have been published [9–11], generally approaching just the alignment step. Siren *et al.* developed a method (GCSA [10]), with construction costs for a whole human genome (plus SNPs) of more than 1 Tb of RAM. Huang *et al.* [11] developed an FM-index [13] encoding of a reference genome-plus-variation ("BWBBLE") by extending the genetic alphabet to encode single-character variants with new characters and then concatenating padded indel variants to the end of the reference genome. We do something similar, but treat all variation in an equivalent manner. While completing this paper, the preprint for GCSA2 was published ([12]), which drops RAM usage of human genome index construction to <100GB at the cost of >1Tb of disk I/O.

We show below how to encode a set of genomes, or a reference plus genetic variation, in an FM-index which naturally distinguishes alternate alleles (defined below). We extend the well known BWT backward search and show how read-mapping can be performed in a way that allows reads to cross multiple variants, allowing recombination to occur naturally. Our data structure supports bidirectional search (which underlies the Super Maximal Exact Match algorithms of bwa-mem [3]), but currently we have only implemented exact matching. We show that index construction for the human genome and 8.3 million genetic variants from the 1000 genomes project uses just 25 GB of RAM on a single core, taking 4.5 hours. We go on to show how inferring a personalised reference results in better genome inference, looking at a highly challenging region - the MSP3.4 gene of *P. falciparum* where alleles differ by one SNP every 3 bp on average.

2   Background: Compressed Text Indexes

Burrows-Wheeler Transform. The Burrows-Wheeler Transform (BWT) of a string is a reversible permutation of its characters that was originally developed for compression [2]. The BWT of a string $T = t_1 t_2 \ldots t_n$ is constructed by sorting its n cyclic shifts $t_1 t_2 \ldots t_n$, $t_2 \ldots t_n t_1$, $\ldots$, $t_n t_1 \ldots t_{n-1}$ in lexicographic order. The matrix obtained is called the Burrows-Wheeler Matrix (BWM) and the sequence from its last column is the BWT. Storing the first and last column of the BWM is sufficient for finding the number of exact matches of a query in T. For locating the position of the matches in T an additional data structure is required, the suffix array.

Suffix Arrays. The suffix array of a string T is an array of integers that provides the starting position of T's suffixes, after they have been ordered lexicographically. Formally, if $T_{i,j}$ is the substring $t_i t_{i+1} \ldots t_j$ of T and SA is the suffix array of T, then $T_{SA[1],n} < T_{SA[2],n} < \cdots < T_{SA[n],n}$. It is related to the BWT, since looking at the substrings preceding the terminating character $ in the BWM rows gives the suffixes of T in lexicographical order.

Backward search Any occurrence of a pattern P in text is a prefix for some suffix of T, so all occurrences will be adjacent in the suffix array of T since suffixes starting with P are sorted together in a SA-interval. Let $C[a]$ be the total number of occurrences in T of characters smaller than a in the alphabet. If P' is a suffix of the query P and $[l(P'), r(P'))$ is its corresponding SA-interval, then the search can be extended to aP' by calculating the new SA-interval:

$$l(aP') = C[a] + rank_a(BWT, l(P') - 1) \tag{1}$$

$$r(aP') = C[a] + rank_a(BWT, r(P)), \tag{2}$$

where the operation $rank_a(S, i)$ returns the number of occurences of symbol a in $S[1, i]$. The search starts with the SA-interval of the empty string, $[1, n]$ and successively adds one character of P in reverse order. When the search is completed, it returns a SA-interval $[l, r)$ for the entire query P. If $r > l$, there are $r - l$ matches for P and their locations in T are given by $SA[i]$ for $l \leq i < r$. Otherwise, the pattern does not exist in T. If the C-array and the ranks have already been stored, the backward search can be performed in $O(|P|)$ time in strings with DNA alphabet.

Wavelet Trees Rank queries scale linearly with the alphabet size by default. The wavelet tree [14] is a data structure designed to store strings with large alphabets efficiently and provide rank calculations in logarithmic time. A string is converted into a balanced binary-tree of bitvectors, whose root is built by taking the sorted alphabet and replacing the lower half of smaller symbols with

a 0, and the other half of larger symbols with a 1. This creates ambiguity initially, but at each tree level, each half of the parent node's alphabet is re-split into 2 and re-encoded, so the ambiguity lessens as the tree is traversed in depth. At the leaves, there is no ambiguity at all. The tree is defined recursively: take the lexicographically ordered alphabet, split it into 2 equal halves; in the string corresponding to the current node (start with original string at root), replace the first half of letters with 0 and the other half with 1; the left child node will contain the 0-encoded symbols and the right child node will contain the 1-encoded symbols, preserving their order from the original string; re-apply the first step for each child node recursively until the alphabet left in each node contains only one or two symbols (so a 0 or 1 determines which symbol it is).
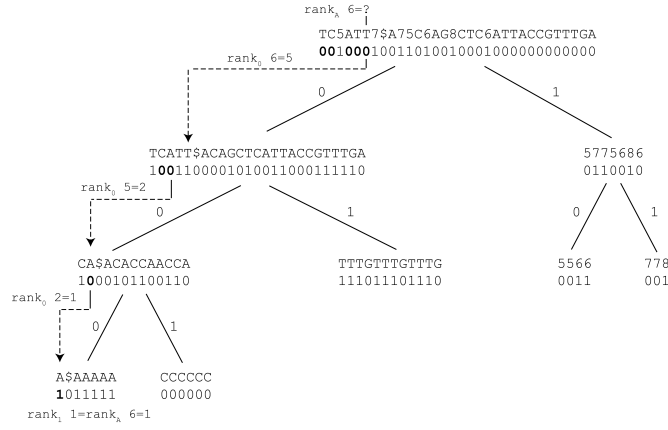
Fig. 1. Wavelet tree encoding of a string. Calculating the rank of the marked "e" is performed by repeated rank() calls moving down the binary tree until the alphabet remaining is just 2 characters. Note that only the bit vectors are stored in the tree, the corresponding strings are only shown here for clarity.

To answer a rank query over the original string with large alphabet, repeated rank queries over the bitvectors in the wavelet tree nodes are used to locate the subtree that contains the leaf where the queried symbol is non-ambiguously encoded. The rank of the queried symbol in this leaf is equal to its rank in the original string. The number of rank queries needed to reach the leaf is equal to the height of the tree, i.e. $\log_2 |\Sigma|$ if we let Σ be the set of symbols in the alphabet. Computing ranks over binary vectors can be done in constant time, so a rank query in a wavelet tree-encoded string has complexity $O(\log_2 |\Sigma|)$.

3   Encoding a variation-aware reference structure

3.1   Terminology

A *variant site* or *site* is a region of the chromosome where there are a number of alternative options for what sequence can be present. These alternatives are termed *alleles* and might be as short as a single character, or could be many hundreds of characters long. A *pan-genome* refers to a representation (with unspecified properties) of a number (greater than 1) of genomes within a species. A Population Reference Graph is an encoding of a pan-genome that enables matching of sequence data to the datastore, inference of nearest mosaic with the appropriate ploidy, and then discovery of new variants not present in the PRG.

3.2   PRG Encoding

Following [8], we use a linear PRG conceptually equivalent to a directed, acyclic, partial order graph, that is generated from a reference sequence and a set of alternative sequences at given variation sites. The graph is linearised into a long string over an alphabet extended with new symbols marking the variants, for which the FM-index can be constructed.

```
Ref  CAAGGCTAT--ACCTACT
Alt1 CAAGGTTATTTACCTGCT    ⟶    CAAGG5CTAT6TTATTT6C5ACCT7A8G7CT
Alt2 CAAGGC-----ACCTACT
```

Fig. 2. A simple PRG linearised according to our encoding. The first site has 3 alleles, which do not here look at all similar, and the second is a SNP.

Building this data structure requires multiple steps.

1. First, corresponding regions of shared (exact match) sequence between the input genomes must be identified. These must be of size k at least (where k is pre-defined), and act as anchors for the coordinates of the variation sites.
2. Second, for any site between two anchor regions, the set of possible haplotypes must be determined from the input genomes, but do not need to be aligned. Indels are naturally supported by haplotypes of different lengths.
3. Each variation site is assigned two unique numeric identifiers, one even and one odd, and they will be called variation markers. The odd identifiers will mark variation site boundaries and will sometimes be referred to as site markers. The even identifiers will mark alternative allele boundaries and will sometimes be referred to as allele boundary markers.
4. For each variation site, its left anchor is added to the linear PRG, followed by its odd identifier. Then each sequence coming from that site, starting with the reference sequence, is successively added to the linear PRG, followed by the even site identifier, except the last sequence, which is followed by the odd identifier.

5. Convert the linear PRG to integer alphabet ($A \to 1$, $C \to 2$, $G \to 3$, $T \to 4$, variation site identifiers $\to$ 5,6,...)
6. The FM-index (suffix array, BWT, wavelet tree over BWT) of the linear PRG is constructed and we will call this the vBWT.

An illustration of these steps on a toy example is given in Figure 2.

Importantly, *the markers force the ends of alternative sequences coming from the same site to be sorted together in a separate block in the Burrows-Wheeler matrix, even if they do not have high sequence similarity.* Therefore, alternative alleles from each site can be queried concurrently.

3.3   Graph structure: constraints

We show in Figure 3a) two sequences which differ by 3 SNPs and give two graph encodings in 3b) and 3c). Both represent the sequence content equally well, and we allow both. In 3d) we have an example where a long deletion lies "over" two other alleles. We would encode this in our PRG as shown in 3e). This works but results in many alternate alleles. An alternative would be to allow "nested" variation, where variants lie on top of other alleles, as shown in Figure 3f). This could be encoded in our system, but we do not allow it for our initial implementation, as it would potentially impact mapping speed.


Fig. 3. PRG graph structure. The sequences shown in Figure 3a) could be represented either as 3 separate mutations (shown in b)), or enumerated as 8 small haplotypes, shown in c). Both are supported by our encoding. Similarly, the sequences in d) could be represented in our implementation as shown in e). However, we do not support "nesting" of alleles, as shown in f).

4   Variation-aware backward search in vBWT

In this section, we present a modified backward search algorithm for exact matching against the vBWT that is aware of alternative sequence paths. When reads align to the non-variable part of the linear PRG or when a variant locus is long

enough to enclose the entire read, the usual backward search algorithm can be used. Otherwise, when the read must cross variation site junctions in order to align, site identifiers and some alternative alleles must be ignored by the search. This means a read can align to multiple substrings of the linear PRG that may not be adjacent in the BWM, so the search can return multiple SA-intervals. We give pseudocode in Algorithm 1 below, and outline the idea in Figure 4.

At each step in the backward search, before extending to the next character, we need to check whether the current matched read substring is preceded by a variation marker anywhere in the linear PRG. A scan for symbols larger than 4 ("range_search_2d" in the pseudocode) must be performed in the BWT within the range given by the current SA-interval. With a wavelet tree this range search can be done in $O(d \log(|\Sigma|/d))$ time, where d is the size of the output. If a variation marker is found and it is an odd number, the read is about to cross a site boundary. The suffix array can be queried to find the position of the two odd numbers (start/end of the site) in the linear PRG.

linear PRG=CAAG**G**5CTAT6**TTATTT**6C5**AC**CT7A8G7CT$

Fig. 4. Backward search across vBWT. We start at the right-hand end of the read GTTATTTAC, with the character C, and as we extend we hit the character 5, signalling the start or end of a variation site. We check the suffix array to get the coordinate in the linear PRG, and find it is the end. Therefore, the read must now continue into one of the alleles, signalled by the number 6. Continuing in this manner (the shorter arrows signify multiple intermediate steps not shown) we are able to align across the site.

If the search cursor is next to the start of the site, it is just the site marker that needs to be skipped, so the SA-interval (size 1) of the suffix starting with that marker needs to be added to the set of intervals that will be extended with the next character in the read. If the search cursor is next to the end of a site, all alternative alleles from that site need to be queried. Their ends are sorted together in the BWM because of the markers, so they can be queried concurrently by adding the SA-interval of suffixes starting with all numbers marking that site (even and odd).

If the variation marker found is an even number, the read is about to cross an allele boundary, which means its current suffix matches the beginning of an alternative allele and the read is about to walk out of a site, so the search cursor needs to jump to the start of site. As previously described, the odd markers corresponding to that site can be found in the sorted first column of the BWM, and then querying the suffix array decides which one marks the start of site. Then the SA-interval (size 1) for the BWM row starting with this odd marker is recorded. Once the check for variation markers is finished and all candidate SA-intervals have been added, each interval can be extended with the next character in the read by using equations 1 and 2. Our implementation leverages the succinct data structures library SDSL [19].

Algorithm 1 Variation-aware backward search

Input: pattern $P[1, m]$ and FM-index of PRG in integer alphabet
Output: list of SA intervals corresponding to matches of P

1: $l \leftarrow C(P[m])$
2: $r \leftarrow C(P[m] + 1)$
3: $i \leftarrow m$
4: `SA_int` $= \{[l, r)\}$ ▷ list of SA intervals
5: `Extra_int` $= \emptyset$ ▷ Extra intervals
6: **while** $i > 0$ and `SA_int` $\neq \emptyset$ **do**
7:     **for all** $[l, r) \in$ `SA_int` **do**
8:         $M \leftarrow$ `WT.range_search_2d`$(l, r - 1, 5, |\Sigma|)$ ▷ find variation site markers
9:         **for all** $(idx, num) \in M$ **do**   ▷ $idx \in [l, r), num \in [5, |\Sigma|]$
10:             **if** $num\%2 = 0$ **then**
11:                 `odd_num` $= num - 1$
12:             **else**
13:                 `odd_num` $= num$
14:             **if** $SA[C(\text{odd\_num})] < SA[C(\text{odd\_num}) + 1]$ **then**
15:                 `start_site` $\leftarrow C(\text{odd\_num})$, `end_site` $\leftarrow C(\text{odd\_num}) + 1$
16:             **else**
17:                 `start_site` $\leftarrow C(\text{odd\_num}) + 1$, `end_site` $\leftarrow C(\text{odd\_num})$
18:             **if** $num\%2 = 1$ and $SA[idx] = SA[\text{site\_end}] + 1$ **then**
19:                 `Extra_int` $=$ `Extra_int` $\cup \{[C(num), C(num + 2))\}$
20:             **else**
21:                 `Extra_int` $=$ `Extra_int` $\cup \{[C[\text{start\_site}], C[\text{start\_site}] + 1)\}$
22:     $i \leftarrow i - 1$
23:     `SA_int` $=$ `SA_int` $\cup$ `Extra_int`
24:     **for all** $[l, r) \in$ `SA_int` **do**
25:         $l = C(P[i]) + rank_{BWT}(P[i], l - 1)$
26:         $r = C(P[i]) + rank_{BWT}(P[i], r)$

---

5 Performance

5.1 Construction cost : the human genome

We constructed a PRG from the human reference genome (GRC37 without "alt" contigs) plus the 1000 genomes final VCF (12GB in size) [6], as described below. We first excluded structural variants which did not have precisely specified alleles, and variants with allele frequency below 5% (rare variation offers limited benefit - our goal is to maximise the proportion of reads mismatching the graph by at most 1 SNP). If two variants occurred at consecutive bases, they were merged into one, and all possible haplotypes enumerated. If the VCF contained two consecutive records which overlapped, the second was discarded. This resulted in a dataset of 7.4 million SNPs and 978000 indels. We give construction costs in Table 1, along with comparative figures for BWBBLE with identical input.

For comparison, using common Finnish SNPs, GCSA took over 1TB of RAM building chromosomes separately and pruning the graph in high diversity regions.

Table 1. FM index construction costs and final data structure size for human reference genome plus 1000 genomes variants

| Software	Peak Memory(GB)	Time (hrs/mins)	Final/in-use memory (GB)
vBWT	25	4h24m	23.5
BWBBLE	60	1h5m	12

GCSA2 reduces the memory footprint to below 128GB RAM, although it prunes the graph in high diversity regions which we want to preserve, running in 13 hours with 32 cores, and using over 1Tb of I/O to fast disk. Our vBWT construction has a lower memory cost than GCSA, GCSA2 and BWBBLE, is faster than GCSA/GCSA2, has no (significant) I/O burden, but is significantly slower than BWBBLE. The memory required for the index after construction is 23.5GB RAM.

### 5.2   Inferring a Closer Reference Genome

Fig. 5. Histogram of number of alleles at each site in MSP3.4 plotted above the chromosome coordinate.)

P. falciparum is a haploid malaria parasite that undergoes recombination within the midgut of vector mosquitoes. It has an unusually repetitive genome (80% AT) that contains more indels than SNPs [15]. There are several regions that present challenges to mapping because samples often diverge strongly from the reference. For example, the merozoite surface protein gene MSP3.4 is known to have two highly diverged lineages at high frequencies in multiple populations from across the world. The lineages differ by around 1 SNP every 3 bases over a 500bp region (the DBL domain) of the gene.

We constructed a catalog of MSP3.4 variation from Cortex [16] and GATK [17] [REMOVE GATK?] variant calls from 650 *P. falciparum* samples from

Fig. 6. Mapping reads from sample 7G8 to *P. falciparum* 3D7 reference genome results in a gap.

Fig. 7. Mapping reads from sample 7G8 to our vBWT-inferred genome removes the gap, leaving isolated variants easy to detect with standard methods

Ghana, Laos and Cambodia, and built a PRG just for that chromosome. We show in Figure 5 the density of variants and number of alleles across the gene. We aligned Illumina 100bp reads from a well-studied sample that was not used in graph construction (named 7G8) to the PRG using backward search (exact matching), and collected counts on the number of reads supporting each allele. At each site we chose the allele with the highest coverage to extract the path through the graph with maximum support - this was our graph-inferred personalised reference for this sample. We then mapped the reads (using bwa_mem [18]) independently to the reference and to the inferred genome. As can be seen in Figures 6 and 7, this gives dramatically better pileup results over the MSP3.4 gene.

5.3   Further performance comments

One consequence of our approach, shared with BWBBLE, is that a single pattern can match multiple intervals in the BWT. Thus, for mapping, we store a list of intervals, along with a corresponding list of which sites and alleles are crossed. Our current implementation of this is naive (using C++ std::vectors and std::lists). We also store an integer array allowing us to determine if a position in the PRG is in a site, and if so, which allele - again naively encoded (in std::vector). For the human example this costs us around 12GB of RAM (not included in table above as not part of index). This array, which contains a zero at every non-variable site in the chromosome, could be stored much more compactly.

This encoding results in an alphabet of size $4 + 2n$ where n is the number of variant sites. We used the wavelet tree to mitigate this cost from $O(|\Sigma|)$ to $O(\log(\Sigma))$. To get an estimate of how this impacts performance in practise, we measure speed of exact matching of 450,000 simulated perfect 150bp reads from each of 100 random haplotypes from *P. falciparum* chromosome 10 with only MSP3.4 variation inserted. To avoid repeatedly calculating the same SA intervals, we precalculate a hash of the SA intervals corresponding to all 9-mers in the PRG. This precalculation was done in 8 minutes using 25 threads. We

compared BWBBLE on the same dataset and, for context, measured speed of traditional mapping of reads with bwa to the haplotype from which reads were sampled. Results (mean) displayed in this table:

Table 2. Exact-matching speed benchmark on *P. falciparum* chromosome 10 with MSP3.4 variation

Software	Time (sec)	Speed (reads/sec)	Unmapped reads
BWA	12	37500	0
vBWT	21	21429	0
BWBBLE	69	6522	?

All experiments were performed on a machine with 64 processors Intel Xeon CPU E5-4620 v2 @ 2.60GHz and 1 TB of memory.

Finally, there is one performance improvement which we have yet to implement - precalculating and storing an array of ranks at marker positions across the BWT - just as in a standard FM-index. This is not normally done for large alphabet wavelet-tree-based suffix arrays, but we can ignore the numeric characters, and store only for A,C,G,T.

6   Discussion

We have demonstrated a whole-genome scale implementation of a PRG designed to enable inference as introduced in [8]. As with any reference graph approach, there is an implicit coupling between mapping and graph structure (for handling alternate alleles). By extending the alphabet and placing positional markers, we are able to ensure that alternate alleles sort together in the BWT matrix, allowing mapping across sites and recombination. In fact, we could encode quite general graph structures in this manner, but for simplicity we imposed a constraint on graph construction, and did not allow nesting of SNPs on other longer alleles. For haploids we naturally infer a personalised reference genome. For other ploidies, our implementation readily lends itself to "lightweight alignment" followed by an HMM to infer a phased tuple of mosaic haplotypes, followed by full MEM-based graph alignment.

Software Our software, gramtools, is available here: http://github.com/iqbal-lab/gramtools, and scripts and PRG files for reproducing this paper are here.

References

1. Valenzuela, D., Valimaki, N., Pitkanen, E., Makinen, V. On enhancing variation detection through pan-genome indexing. Biorxiv. http://dx.doi.org/10.1101/021444
2. Burrows, M., Wheeler, D.J. :A block sorting lossless data compression algorithm. Digital Equipment Corporation, Tech. Rep. 124, 1994.
3. Li, H., Durbin, R.: Fast and accurate short read alignment with Burrows?Wheeler transform. Bioinformatics 25 (14): 1754-1760 (2009)
4. Langmead, B., Salzberg, S.: Fast gapped-read alignment with Bowtie 2. Nature Methods. Mar 4;9(4):357-9 (2012)
5. Reinert, K., Langmead, B., Weese, D., et al.: Alignment of Next-Generation Sequencing Reads. Annu Rev Genomics Hum Genet. 2015;16:133-51
6. The 1000 Genomes Project Consortium: A global reference for human genetic variation. Nature 526, 68?74
7. Ossowski, S., Schneeberger, K., Clark, R.M., et al.. Sequencing of natural strains of Arabidopsis thaliana with short reads. Genome Research 18, 2024-2033 (2008)
8. Dilthey, A., Cox, C., Iqbal, Z., et al.: Improved genome inference in the MHC using a population reference graph. Nature Genetics 47, 682?688 (2015)
9. Schneeberger, K.,Hagmann, J., Ossowski, S., et al.. Simultaneous alignment of short reads against multiple genomes. Genome Biol. 10, R98 (2009).
10. Siren, J., Valimaki, N., Makinen, V. Indexing Graphs for Path Queries with Applications in Genome Research. IEEE/ACM Trans Comput Biol Bioinform. 2014 Mar-Apr;11(2):375-88
11. Huang, L., Popic, V., Batzoglou, S. Short read alignment with populations of genomes. Bioinformatics. Jul 1;29(13):i361-70 (2013)
12. Siren, J. Indexing Variation Graphs. arXiv:1604.06605
13. Ferragina, P. and Manzini, G. Opportunistic data structures with applications. In Proceedings of the 41st Symposium on Foundations of Computer Science (FOCS 2000), IEEE Computer Society, Los Alamitos, CA, USA, pp. 390398 (2000)
14. Grossi, R., Gupta, A. and Vitter, J. High-order entropy-compressed text indexes. In Proceedings of the 14th annual ACM-SIAM symposium on Discrete Algorithms, pages 841850. Society for Industrial and Applied Mathematics, 2003.
15. Miles, A., Iqbal, Z., Vauterin, P., et al. .: Genome variation and meiotic recombination in Plasmodium falciparum: insights from deep sequencing of genetic crosses Biorxiv. http://dx.doi.org/10.1101/024182 (2015)
16. Iqbal, Z., Caccamo, M. Turner, I., et al.: De novo assembly and genotyping of variants using colored de Bruijn graphs. Nature Genetics 44, 226?232 (2012)
17. DePristo, M., Banks, E., Poplin, R.E., et al.: A framework for variation discovery and genotyping using next-generation DNA sequencing data. Nature Genetics 43(5): 491?498 (2011)
18. Li, H. Aligning sequence reads, clone sequences and assembly contigs with BWA-MEM. arXiv:1303.3997
19. Gog, S., Beller, T., Moffat, A. et al.. From Theory to Practice: Plug and Play with Succinct Data Structures. 13th International Symposium on Experimental Algorithms, (SEA 2014) 326-337